

Dexter Fernandes

Bristol • +44 07436007369 • dexterfernandes11@gmail.com • linkedin.com/in/dexterferns •
<https://github.com/Dexter-Fernandes>

Computer Vision Engineer

Computer Vision Engineer with 3+ years of professional engineering experience delivering embedded and cloud-based machine learning solutions. Strong expertise in classical and deep learning-based computer vision with a focus on performance optimisation for edge devices, including Intel Core i7 (VTC7252-7C4IP), Intel Movidius Myriad VPUs, and NVIDIA Jetson AGX Xavier. Proven end-to-end ownership of data pipelines, model development, testing, validation, and production deployment, including hybrid edge-cloud solutions on Microsoft Azure.

WORK EXPERIENCE

Esbaa • Muscat, Oman

Jan 2023 - Aug 2023

AI Engineer

- Engineered and deployed real-time CNN-based perception pipelines on edge devices using PyTorch and OpenVINO, achieving 15% F1 improvement and deterministic 24/7 operation under harsh oil-rig conditions.
- Optimised embedded inference via FP16 quantisation and structured optimisation on Intel VPU hardware, meeting strict latency, memory, and stability constraints for safety-critical use.
- Designed and validated deterministic system behaviour through on-site field testing, performing root-cause analysis of model failures and reducing crash frequency by 67% in live video pipelines.
- Refactored legacy monolith into Kafka-driven, Dockerised microservices, improving fault isolation and increasing MTBF by 35% across edge deployments.
- Owned end-to-end data pipelines, including automated data collection, augmentation, and labelling, increasing dataset diversity by 40% and reducing manual labelling effort by 60%.
- Collaborated with cross-functional engineering and client teams to translate performance metrics into actionable deployment and optimisation decisions.
- Applied Agile and Scrum methodologies to plan sprints, manage deliverables, and ensure iterative software improvement across multidisciplinary teams.

LivNSense Technologies Pvt. Ltd. • Bengaluru, India

May 2020 - Dec 2022

Lead Software Engineer - AI

- Led a 4-member team to design, train, and deploy CNN-based models for real-time object detection, pose estimation, and tracking, reducing false positives by 33% and achieving sub-50 ms inference latency on edge hardware.
- Developed adversarial-condition training and validation conditions simulating occlusion, lighting variation, and motion artefacts, cutting performance variance by 25% and increasing resilience.
- Optimised neural networks using quantisation-aware training, FP16/INT8 fixed-point inference, structured pruning, and layer fusion, reducing model size by 50% while meeting strict memory and latency budgets.
- Owned the full machine learning pipeline from data ingestion and labelling through training, optimisation, deployment, monitoring, and field validation, accelerating iteration cycles and improving deployment reliability.
- Evaluated edge versus cloud inference trade-offs to minimise bandwidth, latency, and operational cost while maintaining safety-critical performance, supporting production cloud-hosted ML workloads on Microsoft Azure.
- Re-architected the embedded inference pipeline and model optimisation workflow, increasing latency from 5 FPS to 15 FPS, enabling real-time detection of safety-critical events.
- Designed data handling and retention practices aligned with GDPR principles, including data minimisation and controlled access, within production computer vision systems.
- Presented key findings and metrics to leadership and clients, translating complex performance results into actionable insights for deployment and further optimisation.

EDUCATION

Master of Science in Robotics - Merit

University of Bristol • Bristol, England

Sep 2023 - Sep 2024

Radiological SLAM with LiDAR Odometry

- Designed and implemented a real-time SLAM system integrating LiDAR and radiological data for the teleoperation of the Clearpath Husky UGV development platform.
- Implemented ROS2 (rclcpp) point cloud pipeline in C++17 using PCL for filtering, segmentation, and registration.
- Developed and maintained ROS2 nodes and interfaces integrating perception outputs into a broader robotic system, managing sensor topics, tf2 frames, and message synchronisation for downstream localisation and navigation.
- Integrated perception pipelines into ROS-based robotic workflows, validating end-to-end data flow from sensors to SLAM and navigation components in Gazebo prior to deployment.
- Simulated and validated 2D/3D mapping and navigation workflows in Gazebo and ROS2 prior to real-world deployment.

Bachelor of Technology in Electrical and Electronics Engineering

Visvesvaraya National Institute of Technology • Nagpur, India

Jul 2016 - May 2020

- Strong foundation in signals, systems, embedded hardware, and control, supporting later work in real-time AI and perception systems.

PROJECTS

ROS2 SLAM & Nav2 Simulation

Oct 2025

- Built a simulated mobile robot integrating 3D LiDAR and RGB camera data in Gazebo Harmonic to perform real-time SLAM and autonomous navigation using ROS 2 Jazzy.
- Applied slam_toolbox and Nav2 for 2D/3D mapping, localisation, and waypoint-based navigation with optimised Ceres backends.
- Processed multi-sensor data streams for spatial alignment and tested mapping accuracy across diverse simulated environments.
- Developed modular ROS 2 launch files and parameterised configurations for flexible experimentation with perception and navigation pipelines.
- Containerised the full system using Docker to ensure reproducibility, version control, and cross-platform deployment.

SKILLS

- **Perception & SLAM:** 3D Feature Extraction, Convolutional Neural Networks, Data Augmentation, Eigen, LiDAR SLAM, NumPy, OpenCV, Point Cloud Library, ROS2 image_transport, Scikit-Image, YOLO
- **Cloud, Edge & MLOps:** Apache Kafka, AWS, Docker, DVC (Data Version Control), Git, Intel VPUs, Linux, Microsoft Azure, MLFlow
- **Programming & Development:** Bash, C++17, CMake, colcon, Pandas, PostgreSQL, Python, rclcpp, rclpy, ROS2 (Humble/Foxy/Jazzy)
- **Deep Learning and Optimisation:** Bayesian Optimisation, DeepSORT, Layer fusion, MultiRocket, ONNX, OpenVINO, Optuna, Post-training Quantisation, PyTorch, Quantisation-Aware Training, Structured pruning, TensorFlow, TensorRT
- **Edge AI and Inference:** Intel Core i7 (VTC7252-7C4IP), Intel Movidius Myriad VPU, NVIDIA Jetson AGX Xavier